

SwiftHire

An automated job search pipeline built for one person — from first alert to cold email in under 10 minutes.

The Problem

Job boards require daily manual checking. By the time most candidates notice a role it's already 3–5 days old, and the best applications come in early. There's no alert system for new postings at specific companies — you either check obsessively or miss things.

Writing a tailored resume for each application takes 30–60 minutes. Most people don't bother — they send the same document everywhere and wonder why response rates are low. The ones who do tailor it spend hours on busywork that doesn't require skill: reading the JD, finding the right bullets to swap in, adjusting language.

Cold outreach gets replies. Applying through the portal alone often doesn't. But finding a recruiter or hiring manager's email takes another 10–20 minutes of searching, and writing a personalised email on top of everything else means most people skip it entirely.

The problem isn't talent — it's that the process is slow enough that most candidates either give up on doing it well or burn out doing it badly.

Step	Time cost (manual)	Notes
Check job boards	Daily, 20–30 min	Repeated across LinkedIn, company sites, ATS portals
Filter irrelevant roles	10–15 min per session	No way to filter by recency, YOE, or sponsorship at once
Read JD and decide	5–10 min	Still required — this stays manual
Tailor resume	30–60 min	Usually skipped, or done poorly under time pressure
Find recruiter contact	10–20 min	LinkedIn, Apollo, guesswork — often unsuccessful
Write cold email	15–30 min	Usually generic. Rarely done alongside an application
Send	2 min	Still requires a person

Who This Is For

SwiftHire was built for one person. The design reflects that — there's no login, no settings page, no multi-user anything. The filters, the company list, the role types, the cities, the notification format — all of it is specific to one job search.

Name	Rohan Jauhari
Degree	MS — Northeastern University, graduating 2026
Target roles	Product Manager (PM), Solutions Engineer (SE), Forward Deployed Engineer (FDE)
Target cities	Boston, NYC, San Francisco, Remote, Chicago, Seattle, Austin, LA, Denver
Company stage	Unicorns, Scale-ups, Enterprise — high-growth tech, AI-native preferred
Constraint	Needs visa sponsorship — roles that explicitly refuse it are filtered out
Experience	Under 4 years — roles requiring 5+ YOE are filtered by default
Timeline	Active search, May 2026

The pipeline is tuned for someone doing a focused, high-volume outreach campaign — not passive job browsing. The goal is to contact relevant people at relevant companies within 48 hours of a role going live, with a resume tailored to that specific job.

The Landscape

Where jobs are posted

The companies worth applying to cluster on two ATS platforms. Greenhouse and Ashby both have public APIs — no scraping, no auth, no rate limits for basic job listing fetches. Lever used to work the same way but their API now returns errors on most public endpoints. Workday requires session cookies and CSRF tokens, making it inaccessible programmatically without a headless browser and regular maintenance.

ATS	API access	Used by	Status
Greenhouse	Public REST API	Anthropic, Stripe, Databricks, Figma, Toast, and most large tech cos	Active
Ashby	Public GraphQL	Ramp, Cursor, Perplexity, Harvey, Sierra — most AI-native startups	Active
Lever	Public API	Some scale-ups and enterprises	Broken
Workday	None	Large enterprises, financial services	Inaccessible
LinkedIn	None	Everything	Blocked

Target companies


23 companies monitored. The list prioritises AI-native companies where FDE and SE roles are growing fastest, unicorns with active hiring, and enterprises with strong PM tracks.

Company	Stage	ATS	Why
Anthropic	Unicorn	Greenhouse	AI safety, strong PM and SE hiring
Cursor	Unicorn	Ashby	AI coding IDE, fastest-growing dev tool
Perplexity	Unicorn	Ashby	AI search, heavy on Solutions and FDE roles
Ramp	Unicorn	Ashby	Finance automation, strong SE and PM culture
Stripe	Unicorn	Greenhouse	Payments infra, PM roles at scale
Databricks	Unicorn	Greenhouse	Data/AI platform, large SE org
Harvey	Scale-up	Ashby	Legal AI, small FDE team growing fast
Sierra	Scale-up	Ashby	Customer AI, early-stage SE and FDE
Figma	Unicorn	Greenhouse	Design platform, strong PM track
Datadog	Enterprise	Greenhouse	Cloud monitoring, large SE org
CoreWeave	Unicorn	Greenhouse	GPU cloud, SE roles for AI infra

Company	Stage	ATS	Why
Brex	Scale-up	Greenhouse	Corporate finance, good PM/SE mix
Cohere	Scale-up	Ashby	Enterprise AI, SE and FDE roles
Cognition	Startup	Ashby	AI agents, small team, early roles
GitLab	Enterprise	Greenhouse	DevSecOps, distributed — good for remote SE
Toast	Enterprise	Greenhouse	Restaurant tech, PM roles with real users
Klaviyo	Scale-up	Greenhouse	Marketing automation, PM and SE tracks
Gusto	Unicorn	Greenhouse	SMB payroll, PM roles with strong scope
Chime	Unicorn	Greenhouse	Neobank, PM and SE growth
Headway	Scale-up	Greenhouse	Mental health, PM roles with impact
Modal	Startup	Ashby	GPU infra, technical SE and FDE
Baseten	Scale-up	Ashby	AI inference, SE and solutions roles
Anyscale	Scale-up	Ashby	AI compute, SE-heavy team

The Pipeline

Five automated steps handle everything between a job posting going live and an email draft landing in Outlook. Three decision points stay with the user.



Step 1 — Continuous monitoring

A Python script runs via launchd every 15 minutes. It calls the Greenhouse and Ashby APIs for each of the 23 companies, classifies results by role type and city, and compares against a record of previously seen job IDs. Anything new fires a macOS notification. Roles older than 3 days don't notify — they're tracked silently. This prevents a flood of alerts whenever a new company or city is added to the list.

◆ **You open the board and decide what's worth pursuing.**

The board filters to the last 2 days, under 4 YOE, and roles where sponsorship isn't explicitly refused. What's left still needs a call — fit, company stage, and whether you're actually interested.

Step 3 — Resume and email, written for the role

Claude has the base resume, full work history, and instructions for how to frame each role type. Paste in the job description and it rewrites the resume around that specific role — pulling out the right projects, adjusting language, reordering sections — and writes a cold email draft: subject line, personalised opening, and body.

Step 4 — Resume download

The Chrome extension watches the Claude tab. When streaming stops it reads the finished LaTeX block and hands it to the browser's native download. The PDF is in Downloads before you've switched tabs.

Step 5 — Contact lookup

One click in the extension popup queries Apollo.io for a recruiter or hiring manager at the target company. Returns a name, title, and email address.

◆ **You confirm it's the right person.**

Wrong contact wastes the entire email regardless of quality.

Step 7 — Email composition

The extension builds an Outlook web compose URL with the confirmed address and what Claude wrote, and opens it in a new tab — recipient, subject, and body already filled in.

◆ **You read the draft, attach the resume, and send.**

Every email that goes out has been read once. Edit anything that doesn't sound right, then send.

What Was Built

Job board — swifthire-board.vercel.app

A React app deployed on Vercel. Static — no backend, no database. The job data is a JavaScript file rewritten every time the poller runs and a new role is found. That also triggers a Vercel redeploy, so the live site stays current automatically.

Filter	Default	Notes
Posted within	2 days	Configurable. Older roles are in the data, just hidden
YOE	Under 4	Toggleable — switch to show all seniority levels
Sponsorship	Hidden	Checked against JD text, not a hardcoded list
City	All	Preset cities + custom add via text input, × to remove
Role type	All	PM, SE, FDE
Company stage	All	Startup, Scale-up, Unicorn, Enterprise
Search	Empty	Matches company, title, description
Sort	Recency	Any column, ascending or descending

Chrome extension — job-to-claude

A Manifest V3 extension with two functions: getting a recruiter contact, and composing the email. No local server. Everything runs in the browser or as a service worker.

Feature	How it works
Resume download	Extension polls the Claude tab. When streaming stops, reads the LaTeX block and triggers chrome.downloads.download()
Contact lookup	Apollo.io query from the popup — returns name, title, email
Email compose	Builds an Outlook web deeplink URL with pre-filled recipient, subject, body. Opens in new tab

Poller — refresh.py

Runs every 15 minutes via launchd. Handles fetching, classifying, deduplication, notifications, sponsorship checks, jobs.js generation, and Vercel redeploy. Logs to /tmp/swifthire-refresh.log.

Decisions and Tradeoffs

Decision	Why	Tradeoff
Static site on Vercel	Zero infrastructure. No server to maintain. Deploys in under 2 min.	New data requires a full redeploy. Page must be reloaded to see changes.
Greenhouse + Ashby only	Both have stable public APIs. Most target companies use one of them.	Misses some companies on Workday or Lever — those require complex scraping.
3-day notification cutoff	Prevents getting alerted for old jobs when adding new companies or cities.	A role posted 4 days ago that's still worth applying to won't notify.
Sponsorship from JD text	More accurate than a hardcoded list. Catches actual policy language.	Only works for Greenhouse. Ashby roles default to noSponsor: false.
No server in the extension	Nothing to maintain, nothing to start before using the extension.	LaTeX compilation can't happen in the browser. Resume export depends on Claude.
Human sends the email	One set of eyes on every email before it goes out.	Adds 2–5 minutes per application.